

# Model selection

In this notebook we're going to analyse different techniques for model selection and afterwards we're going to discuss their shortcomings.

## Selection criteria

First of all, we're going to look at different criteria to compare models based on their performance and complexity.

```
require(ISLR)
```

```
## Loading required package: ISLR
```

```
## Warning: package 'ISLR' was built under R version 3.6.3
```

```
head(Hitters)
```

```
##           AtBat Hits HmRun Runs RBI Walks Years CAtBat CHits CHmRun
## -Andy Allanson    293   66    1  30  29   14    1    293    66    1
## -Alan Ashby       315   81    7  24  38   39   14   3449   835   69
## -Alvin Davis      479  130   18  66  72   76    3   1624   457   63
## -Andre Dawson     496  141   20  65  78   37   11   5628  1575  225
## -Andres Galarra    321   87   10  39  42   30    2    396   101   12
## -Alfredo Griffin  594  169    4  74  51   35   11   4408  1133   19
##           CRuns CRBI CWalks League Division PutOuts Assists Errors
## -Andy Allanson    30   29   14    A         E      446    33    20
## -Alan Ashby       321  414   375    N         W      632    43    10
## -Alvin Davis      224  266   263    A         W      880    82    14
## -Andre Dawson     828  838   354    N         E      200    11    3
## -Andres Galarra    48   46    33    N         E      805    40    4
## -Alfredo Griffin  501  336   194    A         W      282   421   25
##           Salary NewLeague
## -Andy Allanson    NA      A
## -Alan Ashby       475.0    N
## -Alvin Davis      480.0    A
## -Andre Dawson     500.0    N
## -Andres Galarra    91.5     N
## -Alfredo Griffin  750.0    A
```

```
summary(Hitters)
```

```
##           AtBat           Hits           HmRun           Runs
## Min.      : 16.0   Min.       : 1   Min.       : 0.00   Min.       : 0.00
## 1st Qu.:255.2   1st Qu.: 64   1st Qu.: 4.00   1st Qu.: 30.25
## Median :379.5   Median : 96   Median : 8.00   Median : 48.00
## Mean      :380.9   Mean  :101   Mean  :10.77   Mean  : 50.91
## 3rd Qu.:512.0   3rd Qu.:137   3rd Qu.:16.00   3rd Qu.: 69.00
## Max.      :687.0   Max.   :238   Max.   :40.00   Max.   :130.00
##
##           RBI           Walks           Years           CAtBat
```

```
## Min. : 0.00 Min. : 0.00 Min. : 1.000 Min. : 19.0
## 1st Qu.: 28.00 1st Qu.: 22.00 1st Qu.: 4.000 1st Qu.: 816.8
## Median : 44.00 Median : 35.00 Median : 6.000 Median : 1928.0
## Mean : 48.03 Mean : 38.74 Mean : 7.444 Mean : 2648.7
## 3rd Qu.: 64.75 3rd Qu.: 53.00 3rd Qu.: 11.000 3rd Qu.: 3924.2
## Max. : 121.00 Max. : 105.00 Max. : 24.000 Max. : 14053.0
##
## CHits CHmRun CRuns CRBI
## Min. : 4.0 Min. : 0.00 Min. : 1.0 Min. : 0.00
## 1st Qu.: 209.0 1st Qu.: 14.00 1st Qu.: 100.2 1st Qu.: 88.75
## Median : 508.0 Median : 37.50 Median : 247.0 Median : 220.50
## Mean : 717.6 Mean : 69.49 Mean : 358.8 Mean : 330.12
## 3rd Qu.: 1059.2 3rd Qu.: 90.00 3rd Qu.: 526.2 3rd Qu.: 426.25
## Max. : 4256.0 Max. : 548.00 Max. : 2165.0 Max. : 1659.00
##
## CWalks League Division PutOuts Assists
## Min. : 0.00 A:175 E:157 Min. : 0.0 Min. : 0.0
## 1st Qu.: 67.25 N:147 W:165 1st Qu.: 109.2 1st Qu.: 7.0
## Median : 170.50 Median : 212.0 Median : 39.5
## Mean : 260.24 Mean : 288.9 Mean : 106.9
## 3rd Qu.: 339.25 3rd Qu.: 325.0 3rd Qu.: 166.0
## Max. : 1566.00 Max. : 1378.0 Max. : 492.0
##
## Errors Salary NewLeague
## Min. : 0.00 Min. : 67.5 A:176
## 1st Qu.: 3.00 1st Qu.: 190.0 N:146
## Median : 6.00 Median : 425.0
## Mean : 8.04 Mean : 535.9
## 3rd Qu.: 11.00 3rd Qu.: 750.0
## Max. : 32.00 Max. : 2460.0
## NA's : 59
```

```
# removing the NA
```

```
dim(Hitters)
```

```
## [1] 322 20
```

```
Hitters<- na.omit(Hitters)
```

```
dim(Hitters)
```

```
## [1] 263 20
```

We're going to use cross-validation to compare the results from different selection criteria.

```
nfolds <- 10
```

```
n <- dim(Hitters)[1]
```

```
folds <- cut(1:n, nfolds, labels = F)
```

```
# a bit of shuffling
```

```
indices <- sample(1:n, size=n, replace=F)
```

```
library(leaps)
```

```
## Warning: package 'leaps' was built under R version 3.6.3
```

```
get.bss.test.error<- function(train, test, cv.best){
```

```
  # estimates the error on the test dataset for the best model
```

```
  # according to each criteria
```

```
  all.best<- regsubsets(x=Salary~.,data=train,nbest=1,
```

```

        nvmax=dim(train)[2]-1, # using all variables
        method="forward" )
s <- summary(all.best)
r2 <- coef(all.best, id=which.max(s$rsq))
adjr2 <- coef(all.best, id=which.max(s$adjr2))
cp <- coef(all.best, id=which.min(s$cp))
bic <- coef(all.best, id=which.min(s$bic))
cv.coefs <- coef(all.best, id=cv.best)
# test predictions
r2.pred <- model.matrix(Salary~.,test)[,names(r2)]%*%r2
adjr2.pred <- model.matrix(Salary~.,test)[,names(adjr2)]%*%adjr2
cp.pred <- model.matrix(Salary~.,test)[,names(cp)]%*%cp
bic.pred <- model.matrix(Salary~.,test)[,names(bic)]%*%bic
cv.pred <- model.matrix(Salary~.,test)[,names(cv.coefs)]%*%cv.coefs
# test errors
errors <- mean((r2.pred - test$Salary)**2)
errors <- c(errors,mean((adjr2.pred - test$Salary)**2))
errors <- c(errors,mean((cp.pred - test$Salary)**2))
errors <- c(errors,mean((bic.pred - test$Salary)**2))
errors <- c(errors,mean((cv.pred - test$Salary)**2))
return(errors)
}

get.cv.error <- function(ncv, nmodels, data){
  # evaluates the mean cross-validation error of the linear model
  # with the selected coefficients
  n.cv <- dim(data)[1]
  folds.cv <- cut(1:n.cv, ncv, labels=F)
  cv.errors <- matrix(nrow = ncv, ncol = nmodels)
  indices.cv <- 1:n.cv
  for(j in 1:ncv){
    test.indices.cv <- indices.cv[folds.cv==j]
    test.cv <- data[test.indices.cv,]
    train.cv <- data[-test.indices.cv,]
    cv.all.best<- regsubsets(x=Salary~.,data=train.cv,
                           nbest=1,nvmax=nmodels, # using all variables
                           method="forward" )

    for(m in 1:nmodels){
      cv.coefs <- coef(cv.all.best, id=m)
      cv.preds <- model.matrix(Salary~.,test)[,names(cv.coefs)]%*%cv.coefs
      # test errors
      cv.errors[j,m] <- mean((cv.preds - test$Salary)**2)
    }
  }
  # selecting the model with the least mean error
  # expected test MSE estimated by CV for each model
  return(which.min(colMeans(cv.errors)))
}

test.errors <- matrix(nrow=nfolds, ncol=5)

for(i in 1:nfolds){

```

```

test.indices <- indices[folds==i]
test <- Hitters[test.indices,]
train <- Hitters[-test.indices,]
# Now we'll use BSS on the train dataset
# And we'll record the error on the test set
# get best cv model
cv.best <- get.cv.error(ncv=5, nmodels=(dim(Hitters)[2]-1),data = train)
test.errors[i,] <- get.bss.test.error(train=train, test=test, cv.best=cv.best)
}

```

Let's look at the results.

```

test.errors <- data.frame(test.errors)
names(test.errors) <- c("r2", "adjr2", "cp", "bic", "cv")
test.errors

```


##	r2	adjr2	cp	bic	cv
## 1	160378.72	143422.40	143422.40	159610.15	143422.40
## 2	80984.61	80957.27	83288.87	95476.85	79866.42
## 3	199024.56	204565.73	204565.73	205284.86	193625.86
## 4	88793.36	86079.44	86995.54	85359.33	85359.33
## 5	150460.62	141612.89	134504.21	137837.55	142262.61
## 6	114399.12	102855.15	100042.77	89121.79	97775.38
## 7	198900.94	208226.35	181812.81	168333.41	177456.26
## 8	95930.56	93306.29	93306.29	75335.05	86151.64
## 9	82647.53	81797.56	83679.26	93913.87	83031.22
## 10	56376.12	60188.58	60188.58	73283.04	53540.29

```

plot(1:10, test.errors$r2, type="l", lty="dashed", col=2, ylab="test error", main="cv MSE estimate ", lwd=2)
lines(1:10, test.errors$adjr2, type="l", lty="dashed", col=3, lwd=2)
lines(1:10, test.errors$cp, type="l", lty="dashed", col=4, lwd=2)
lines(1:10, test.errors$bic, type="l", lty="dashed", col=5, lwd=2)
lines(1:10, test.errors$cv, type="l", lty="dashed", col=6, lwd=2)
legend("topright", legend = c("r2", "adjr2", "cp", "bic", "cv"), col=c(2,3,4,5,6), lty="dashed")

```

## cv MSE estimate



```
colMeans(test.errors)
```

```
##      r2      adjr2      cp      bic      cv
## 122789.6 120301.2 117180.6 118355.6 114249.1
```

```
which.min(colMeans(test.errors))
```

```
## cv
## 5
```

So the cross validation criteria seems to be the most reliable in model selection. We'll now use this criteria to select the best model fitting it on the whole data.

```
best.cv <- get.cv.error(ncv=10, nmodels=(dim(Hitters)[2]-1), data=Hitters)
best.cv
```

```
## [1] 2
```

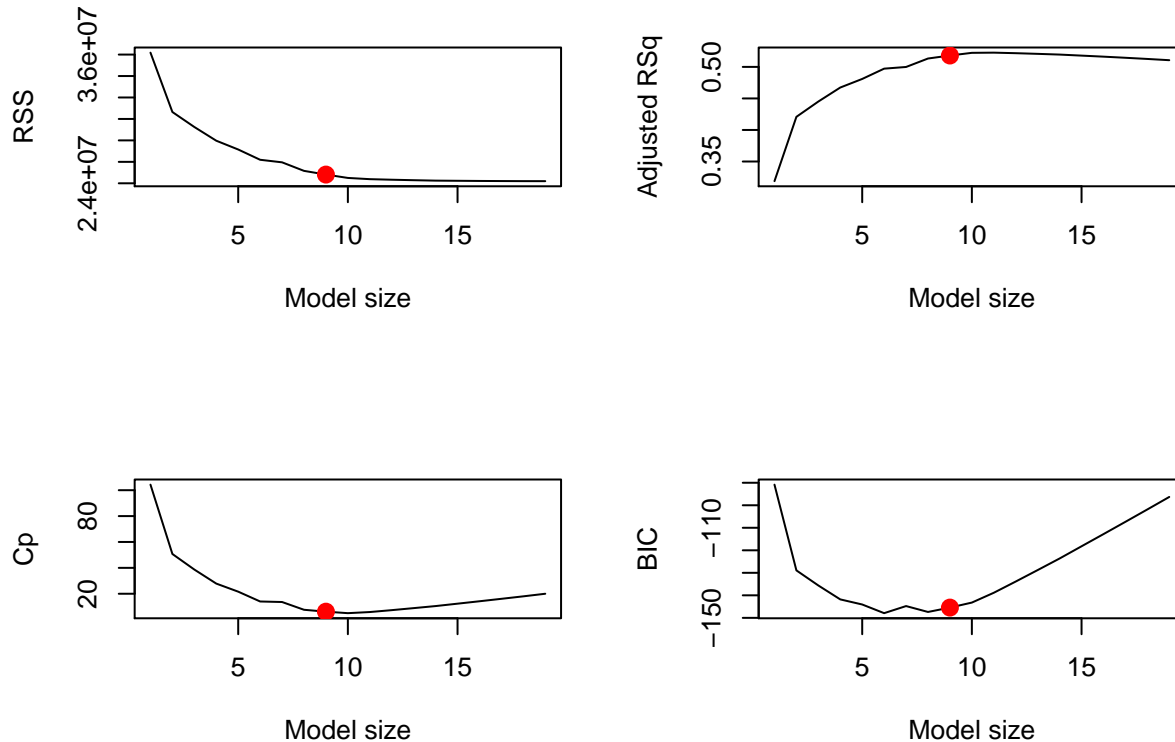
Let's now look at the best model:

```
nmodels <- 19
all.best <- regsubsets(x=Salary~.,data=Hitters,
                      nbest=1,nvmax=nmodels, # using all variables
                      method="forward")
coef(all.best, id=9)
```


```
##      (Intercept)      AtBat      Hits      Walks      CAtBat
## 146.24960033    -1.93676754    6.65672102    5.55204413   -0.09953904
##           CRuns           CRBI           CWalks      DivisionW      PutOuts
## 1.25067124    0.66176849   -0.77798498  -115.34950146    0.27773062
```

Let's look at some plots to see how the best model is seen according to other criteria.


```
par(mfrow=c(2,2))
s <- summary(all.best)
# rss
plot(s$rss,xlab="Model size",ylab="RSS",type="l")
points(9, s$rss[9], col="red",cex=2,pch=20)
# adjr2
plot(s$adjr2,xlab="Model size",ylab="Adjusted RSq",type="l")
points(9, s$adjr2[9], col="red",cex=2,pch=20)
# Cp
plot(s$cp,xlab="Model size",ylab="Cp",type='l')
points(9, s$cp[9], col="red", cex=2, pch=20)
# BIC
plot(s$bic,xlab="Model size",ylab="BIC",type='l')
points(9, s$bic[9], col="red", cex=2, pch=20)
```



```
par(mfrow=c(1,1))
plot(all.best, scale="r2")
```




```
plot(all.best, scale="adjr2")
```




```
plot(all.best, scale="Cp")
```





```
plot(all.best, scale="bic")
```



## Shrinkage methods

We'll now look at a different set of selection tools: shrinkage methods like LASSO and RIDGE. We'll actually use the ElasticNet model, of which Lasso and Ridge are special cases.

```
require(hdi)
```

```
## Loading required package: hdi
## Warning: package 'hdi' was built under R version 3.6.3
## Loading required package: scalreg
## Loading required package: lars
## Loaded lars 1.2
```

```
data("riboflavin")
```

The riboflavin dataset records the riboflavin production by *Bacillus subtilis* together with the gene expressions. Each row refers to one gene, storing in *x* its expression level.

```
require(glmnet)
```

```
## Loading required package: glmnet
## Warning: package 'glmnet' was built under R version 3.6.3
## Loading required package: Matrix
## Loaded glmnet 3.0-2
```

```
attach(riboflavin)
lambda.grid <- grid <- 10^seq(10,-2, length = 100)
lasso <- glmnet(x = x, y=y, alpha = 1, lambda = lambda.grid)
ridge <- glmnet(x = x, y=y, alpha = 0, lambda = lambda.grid)
```

Exploration of the output:

```
dim(coef(lasso))
```

```
## [1] 4089 100
```

```
dim(coef(ridge))
```

```
## [1] 4089 100
```

```
lasso$lambda[50]
```

```
## [1] 11497.57
```

```
round(coef(lasso)[,50],2)[coef(lasso)[,50]!=0]
```

```
## (Intercept)
```

```
## -7.16
```

```
ridge$lambda[50]
```

```
## [1] 11497.57
```

```
round(coef(ridge)[,50],2)[round(coef(ridge)[,50],2)!=0]
```


```
## (Intercept)
```

```
## -7.15
```


The above lambda is too high to allow any value to be different from zero. Let's have a broader look at the results with a plot:

```
plot(lasso)
```

```
## Warning in regularize.values(x, y, ties, missing(ties)): collapsing to unique
## 'x' values
```




```
plot(ridge)
```



The above plots are a perfect synthesis of the differences between ridge and lasso: while ridge performs a “soft thresholding”, slowly shrinking all variables to 0 as lambda increases, lasso performs a “hard thresholding”, cutting off variables as they reach a certain threshold (determined by lambda).

```
plot(ridge, xvar="lambda")
```



How to use a model for predictions?

```
preds <- predict(lasso, s = 0.01, newx=riboflavin$x)
mse.train <- mean((riboflavin$y - preds)**2)
mse.train
```


```
## [1] 0.005253187
```

```
preds.lasso <- predict(lasso, s = grid, newx=riboflavin$x)
plot(log(grid), colMeans((riboflavin$y - preds.lasso)**2), col="red", main="Train MSE", type="l", ylab = "Train MSE")

preds.ridge <- predict(ridge, s = grid, newx=riboflavin$x)
lines(log(grid), colMeans((riboflavin$y - preds.ridge)**2), col="blue", type="l", ylab = "", lty="dashed")

legend("right", legend=c("lasso","ridge"), col=c("red","blue"), lty="dashed")
```

## Train MSE



Now let's use cross validation to do both model assessment and model selection: we're going to select the best lasso and ridge models on a train dataset and evaluate them with a second cross-validation against hold-out sets.

```
nfolds <- 10
n <- dim(riboflavin)[1]
folds <- cut(1:n, breaks = nfolds, labels = F)
#shuffling
indices <- sample(1:n, size=n, replace = F)
res.cv <- matrix(nrow=10, ncol=4)
for(i in 1:nfolds){
  test.indices <- indices[folds==i]
  test <- riboflavin[test.indices, ]
  train <- riboflavin[-test.indices, ]
  ## model selection using cv on the train dataset
  ## LASSO
  lasso.cv <- cv.glmnet(x=train$x, y=train$y,
                       alpha = 1, lambda = lambda.grid, nfolds = 10)
  lasso.lambda <- lasso.cv$lambda.min
  res.cv[i,1] <- lasso.lambda
  lasso.fit <- glmnet(train$x, train$y, alpha = 1, lambda = lasso.lambda)
  ## RIDGE
  ridge.cv <- cv.glmnet(x=train$x, y=train$y,
                       alpha = 0, lambda = lambda.grid, nfolds = 10)
  ridge.lambda <- ridge.cv$lambda.min
  res.cv[i,3] <- ridge.lambda
  ridge.fit <- glmnet(train$x, train$y, alpha = 0, lambda = ridge.lambda)
```

```



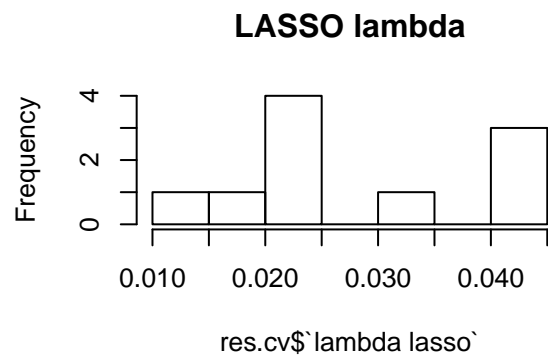

## model assessment on the test dataset
## LASSO
lasso.predict <- predict(lasso.fit, newx=test$x)
lasso.error<- mean((test$y-lasso.predict)**2)
res.cv[i,2]<-lasso.error
## RIDGE
ridge.predict <- predict(ridge.fit, newx=test$x)
ridge.error<- mean((test$y-ridge.predict)**2)
res.cv[i,4]<-ridge.error
}
res.cv <- data.frame(res.cv)
names(res.cv) <- c("lambda lasso","error lasso","lambda ridge", "error ridge")
res.cv

##      lambda lasso error lasso lambda ridge error ridge
## 1      0.02310130  0.2283152    6.1359073  0.23587021
## 2      0.04037017  0.5763044    6.1359073  0.55236718
## 3      0.04037017  0.3959881    8.1113083  0.39342692
## 4      0.01000000  0.2913871    1.5199111  0.22009343
## 5      0.01747528  0.1397455    6.1359073  0.07637653
## 6      0.03053856  0.1279141    6.1359073  0.11294895
## 7      0.02310130  0.1406088    0.6579332  0.36587483
## 8      0.02310130  0.3240908    3.5111917  0.31441988
## 9      0.04037017  0.1396851    6.1359073  0.13475449
## 10     0.02310130  0.1864312    3.5111917  0.23379612

par(mfrow=c(2,2))
hist(res.cv$error lasso`, main="LASSO error")
hist(res.cv$lambda lasso`, main="LASSO lambda")
hist(res.cv$error ridge`, main="RIDGE error")
hist(res.cv$lambda ridge`, main="RIDGE lamdba")


```





Not only the ridge presents a higher expected error on the test set, but it also has more variability in the lambda values. Therefore we're going to use Lasso, fitting it to the whole dataset.

```
lasso.cv <- cv.glmnet(x=x, y=y, alpha = 1, lambda = lambda.grid, nfolds =10)
plot(lasso.cv)
```



```
best.lambda <- lasso.cv$lambda.min
lasso.fit <- glmnet(x,y, alpha=1, lambda=best.lambda, thresh=1e-12)
sum(coef(lasso.fit)!=0)
```

```
## [1] 41
```

In the final model only 41 of the initial >4000 genes are active.